

Data Preparation and Exploration

1. Load and Explore a Dataset:

Load a dataset (e.g., Iris, Wine, Breast Cancer) and display the first few rows.

```
In [81]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [109... df= pd.read_csv("iris.data.csv")
df
```

```
Out[109]:
```

	5.1	3.5	1.4	0.2	Iris-setosa
0	4.9	3.0	1.4	0.2	Iris-setosa
1	4.7	3.2	1.3	0.2	Iris-setosa
2	4.6	3.1	1.5	0.2	Iris-setosa
3	5.0	3.6	1.4	0.2	Iris-setosa
4	5.4	3.9	1.7	0.4	Iris-setosa
...
144	6.7	3.0	5.2	2.3	Iris-virginica
145	6.3	2.5	5.0	1.9	Iris-virginica
146	6.5	3.0	5.2	2.0	Iris-virginica
147	6.2	3.4	5.4	2.3	Iris-virginica
148	5.9	3.0	5.1	1.8	Iris-virginica

149 rows × 5 columns

2. Inspect the Data:

Display the data types of each column and summary statistics.

```
In [110... df.columns
```

```
Out[110]: Index(['5.1', '3.5', '1.4', '0.2', 'Iris-setosa'], dtype='object')
```

```
In [111... df.describe()
```

```
Out[111]:
```

	5.1	3.5	1.4	0.2
count	149.000000	149.000000	149.000000	149.000000
mean	5.848322	3.051007	3.774497	1.205369
std	0.828594	0.433499	1.759651	0.761292
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.400000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```
In [112...] df.columns #find columns names
```

```
Out[112]: Index(['5.1', '3.5', '1.4', '0.2', 'Iris-setosa'], dtype='object')
```

```
In [128...] df.shape
```

```
Out[128]: (149, 5)
```

```
In [113...] x = df.iloc[:, :-1] # Selects all rows and all columns except the last one
y = df.iloc[:, -1] # Selects all rows and the last column
```

```
In [114...] x
```

```
Out[114]:
```

	5.1	3.5	1.4	0.2
0	4.9	3.0	1.4	0.2
1	4.7	3.2	1.3	0.2
2	4.6	3.1	1.5	0.2
3	5.0	3.6	1.4	0.2
4	5.4	3.9	1.7	0.4
...
144	6.7	3.0	5.2	2.3
145	6.3	2.5	5.0	1.9
146	6.5	3.0	5.2	2.0
147	6.2	3.4	5.4	2.3
148	5.9	3.0	5.1	1.8

149 rows × 4 columns

```
In [115...] print(y)
```

```

0      Iris-setosa
1      Iris-setosa
2      Iris-setosa
3      Iris-setosa
4      Iris-setosa
...
144    Iris-virginica
145    Iris-virginica
146    Iris-virginica
147    Iris-virginica
148    Iris-virginica
Name: Iris-setosa, Length: 149, dtype: object

```

In [116...]

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 149 entries, 0 to 148
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    5.1         149 non-null    float64
1    3.5         149 non-null    float64
2    1.4         149 non-null    float64
3    0.2         149 non-null    float64
4    Iris-setosa 149 non-null    object
dtypes: float64(4), object(1)
memory usage: 5.9+ KB

```

know the target variable

In [120...]

```

print(df['Iris-setosa'].unique())

['Iris-setosa' 'Iris-versicolor' 'Iris-virginica']

```

In [121...]

```

df.dtypes

```

Out[121]:

```

5.1      float64
3.5      float64
1.4      float64
0.2      float64
Iris-setosa  object
dtype: object

```

3. Check for Missing Values:

Identify and handle missing values in the dataset.

In [122...]

```

df.isnull().sum()

```

Out[122]:

```

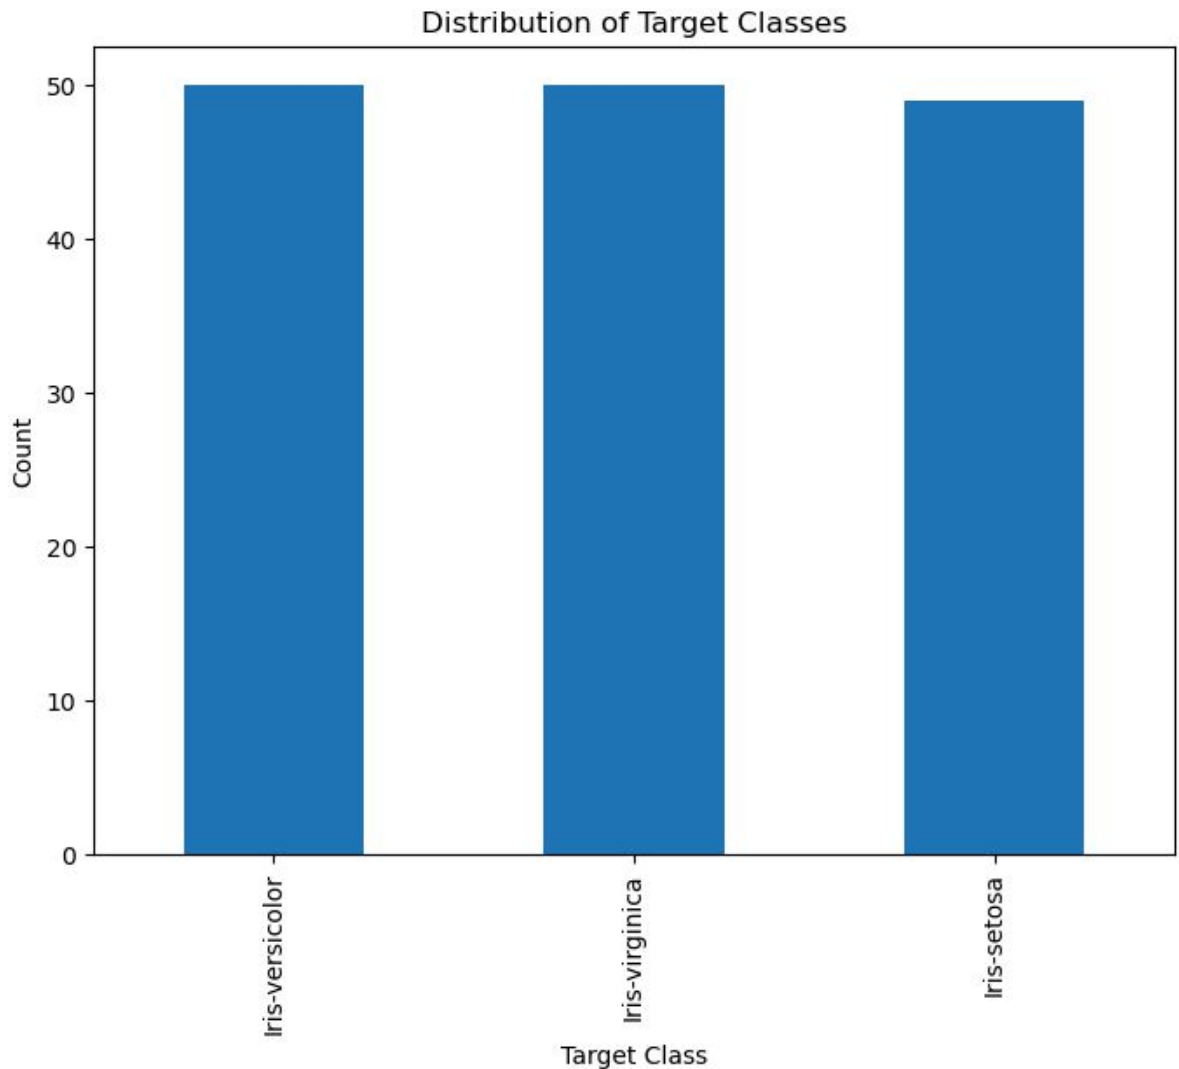
5.1      0
3.5      0
1.4      0
0.2      0
Iris-setosa  0
dtype: int64

```

4. Data Visualization:

Visualize the distribution of classes in the target variable

```
In [124... plt.figure(figsize=(8, 6))
df['Iris-setosa'].value_counts().plot(kind='bar')
plt.title('Distribution of Target Classes')
plt.xlabel('Target Class')
plt.ylabel('Count')
plt.show()
```



5. Feature Scaling:

Scale the features using StandardScaler or MinMaxScaler.

```
In [125... from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
In [126... scaler = StandardScaler()
X_scaled = scaler.fit_transform(x)
```

6. Train a Logistic Regression Model:

Train and evaluate a logistic regression model.

```
In [129... x_train, x_test, y_train, y_test= train_test_split(X_scaled, y, test_size=0.2, rand

In [130... from sklearn.linear_model import LogisticRegression

In [131... lr= LogisticRegression()
lr.fit(x_train, y_train)
y_pred= lr.predict(x_test)

In [133... from sklearn.metrics import accuracy_score

In [134... print(f'Logistic Regression Accuracy: {accuracy_score(y_test, y_pred)}')

Logistic Regression Accuracy: 0.9
```

7. Train a K-Nearest Neighbors (KNN) Model:

Train and evaluate a KNN model.

```
In [135... from sklearn.neighbors import KNeighborsClassifier

In [137... knn = KNeighborsClassifier()
knn.fit(x_train, y_train)
y_pred_knn = knn.predict(x_test)
print(f'KNN Accuracy: {accuracy_score(y_test, y_pred_knn)}')

KNN Accuracy: 0.9333333333333333

C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:22
8: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.
11.0, this behavior will change: the default value of `keepdims` will become Fals
e, the `axis` over which the statistic is taken will be eliminated, and the value
None will no longer be accepted. Set `keepdims` to True or False to avoid this war
ning.
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

8. Train a Decision Tree Model:

Train and evaluate a decision tree model.

```
In [138... from sklearn.tree import DecisionTreeClassifier

In [139... dt = DecisionTreeClassifier()
dt.fit(x_train, y_train)
y_pred_dt = dt.predict(x_test)
print(f'Decision Tree Accuracy: {accuracy_score(y_test, y_pred_dt)}')

Decision Tree Accuracy: 0.9
```

9. Train a Random Forest Model:

Train and evaluate a random forest model.

```
In [140... from sklearn.ensemble import RandomForestClassifier
```

```
In [141... rf = RandomForestClassifier()  
rf.fit(x_train, y_train)  
y_pred_rf = rf.predict(x_test)  
print(f'Random Forest Accuracy: {accuracy_score(y_test, y_pred_rf)}')
```

Random Forest Accuracy: 0.9

10. Train a Support Vector Machine (SVM) Model:

Train and evaluate an SVM model.

```
In [142... from sklearn.svm import SVC
```

```
In [143... svm = SVC()  
svm.fit(x_train, y_train)  
y_pred_svm = svm.predict(x_test)  
print(f'SVM Accuracy: {accuracy_score(y_test, y_pred_svm)}')
```

SVM Accuracy: 0.9333333333333333

11. Train a Naive Bayes Model:

Train and evaluate a naive bayes model.

```
In [144... from sklearn.naive_bayes import GaussianNB
```

```
In [145... nb = GaussianNB()  
nb.fit(x_train, y_train)  
y_pred_nb = nb.predict(x_test)  
print(f'Naive Bayes Accuracy: {accuracy_score(y_test, y_pred_nb)}')
```

Naive Bayes Accuracy: 0.8666666666666667

12. Train a Gradient Boosting Model:

Train and evaluate a gradient boosting model.

```
In [146... from sklearn.ensemble import GradientBoostingClassifier
```

```
In [147... gb = GradientBoostingClassifier()  
gb.fit(x_train, y_train)  
y_pred_gb = gb.predict(x_test)  
print(f'Gradient Boosting Accuracy: {accuracy_score(y_test, y_pred_gb)}')
```

Gradient Boosting Accuracy: 0.9

13. Train a LightGBM Model:

Train and evaluate a LightGBM model

In [151...

```
pip install lightgbm
```

```
Collecting lightgbm
  Downloading lightgbm-4.5.0-py3-none-win_amd64.whl (1.4 MB)
    ----- 1.4/1.4 MB 6.5 MB/s eta 0:00:00
Requirement already satisfied: numpy>=1.17.0 in c:\users\shaw3\anaconda3\lib\site-packages (from lightgbm) (1.24.4)
Requirement already satisfied: scipy in c:\users\shaw3\anaconda3\lib\site-packages (from lightgbm) (1.9.1)
Installing collected packages: lightgbm
Successfully installed lightgbm-4.5.0
Note: you may need to restart the kernel to use updated packages.
```

In [153...

```
import lightgbm as lgb
```

In [154...

```
lgb_model = lgb.LGBMClassifier()
lgb_model.fit(x_train, y_train)
y_pred_lgb = lgb_model.predict(x_test)
print(f'LightGBM Accuracy: {accuracy_score(y_test, y_pred_lgb)}')
```


[illegible]

[illegible]

[illegible]

Perform hyperparameter tuning for logistic regression using GridSearchCV.

```
In [156... from sklearn.model_selection import GridSearchCV
```

```
In [ ]: # C: This is the inverse of regularization strength.  
        #Smaller values specify stronger regularization. The grid specifies three possi
```

```
In [ ]: #penalty: This specifies the type of regularization to use. 'L1' corresponds to Lasso
        #and 'L2' corresponds to Ridge (L2) regularization
```

```
In [157... lr_params = {'C': [0.1, 1, 10], 'penalty': ['l1', 'l2']}
lr_grid = GridSearchCV(LogisticRegression(), lr_params, cv=5)
lr_grid.fit(x_train, y_train)
print(f'Best Logistic Regression Params: {lr_grid.best_params_}')
print(f'Logistic Regression Grid Search CV Score: {lr_grid.best_score_}')
```

```
Best Logistic Regression Params: {'C': 10, 'penalty': 'l2'}
Logistic Regression Grid Search CV Score: 0.9829710144927537
```

```
C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\model_selection\_validation.py:
372: FitFailedWarning:
15 fits failed out of a total of 30.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_sco
re='raise'.
```

Below are more details about the failures:

15 fits failed with the following error:

Traceback (most recent call last):

```
File "C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\model_selection\_validation.py", line 680, in _fit_and_score
```

```
estimator.fit(X_train, y_train, **fit_params)
```

```
File "C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py", line 1461, in fit
```

```
solver = _check_solver(self.solver, self.penalty, self.dual)
```

```
File "C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py", line 447, in _check_solver
```

```
raise ValueError(
```

```
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.
```

```
warnings.warn(some fits failed message, FitFailedWarning)
```

```
C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\model_selection\search.py:969:
```

```
UserWarning: One or more of the test scores are non-finite: [      nan  0.90833333
 nan  0.96666667      nan  0.98297101]
```

```
warnings.warn(
```

Hyperparameter Tuning for KNN:

Perform hyperparameter tuning for KNN using GridSearchCV.

```
In [158... knn_params = {'n_neighbors': [3, 5, 7], 'weights': ['uniform', 'distance']}
knn_grid = GridSearchCV(KNeighborsClassifier(), knn_params, cv=5)
knn_grid.fit(x_train, y_train)
print(f'Best KNN Params: {knn_grid.best_params_}')
print(f'KNN Grid Search CV Score: {knn_grid.best_score}')
```

```
Best KNN Params: {'n_neighbors': 5, 'weights': 'uniform'}
KNN Grid Search CV Score: 0.9663043478260869
```

C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```



```

C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:22
8: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.
11.0, this behavior will change: the default value of `keepdims` will become Fals
e, the `axis` over which the statistic is taken will be eliminated, and the value
None will no longer be accepted. Set `keepdims` to True or False to avoid this war
ning.
    mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:22
8: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.
11.0, this behavior will change: the default value of `keepdims` will become Fals
e, the `axis` over which the statistic is taken will be eliminated, and the value
None will no longer be accepted. Set `keepdims` to True or False to avoid this war
ning.
    mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:22
8: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.
11.0, this behavior will change: the default value of `keepdims` will become Fals
e, the `axis` over which the statistic is taken will be eliminated, and the value
None will no longer be accepted. Set `keepdims` to True or False to avoid this war
ning.
    mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:22
8: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.
11.0, this behavior will change: the default value of `keepdims` will become Fals
e, the `axis` over which the statistic is taken will be eliminated, and the value
None will no longer be accepted. Set `keepdims` to True or False to avoid this war
ning.
    mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:22
8: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.
11.0, this behavior will change: the default value of `keepdims` will become Fals
e, the `axis` over which the statistic is taken will be eliminated, and the value
None will no longer be accepted. Set `keepdims` to True or False to avoid this war
ning.
    mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:22
8: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.
11.0, this behavior will change: the default value of `keepdims` will become Fals
e, the `axis` over which the statistic is taken will be eliminated, and the value
None will no longer be accepted. Set `keepdims` to True or False to avoid this war
ning.
    mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
C:\Users\shaw3\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:22
8: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the
default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.
11.0, this behavior will change: the default value of `keepdims` will become Fals
e, the `axis` over which the statistic is taken will be eliminated, and the value
None will no longer be accepted. Set `keepdims` to True or False to avoid this war
ning.
    mode, _ = stats.mode(_y[neigh_ind, k], axis=1)

```

Hyperparameter Tuning for Decision Tree:

Perform hyperparameter tuning for decision tree using GridSearchCV


```
In [160... dt_params = {'max_depth': [2, 4, 6], 'min_samples_split': [2, 4, 6]}
dt_grid = GridSearchCV(DecisionTreeClassifier(), dt_params, cv=5)
dt_grid.fit(x_train, y_train)
print(f'Best Decision Tree Params: {dt_grid.best_params_}')
print(f'Decision Tree Grid Search CV Score: {dt_grid.best_score_}')
```

Best Decision Tree Params: {'max_depth': 6, 'min_samples_split': 2}
Decision Tree Grid Search CV Score: 0.9916666666666668

Hyperparameter Tuning for Random Forest:

Perform hyperparameter tuning for random forest using GridSearchCV.

```
In [161... rf_params = {'n_estimators': [50, 100, 150], 'max_depth': [2, 4, 6]}
rf_grid = GridSearchCV(RandomForestClassifier(), rf_params, cv=5)
rf_grid.fit(x_train, y_train)
print(f'Best Random Forest Params: {rf_grid.best_params_}')
print(f'Random Forest Grid Search CV Score: {rf_grid.best_score_}')
```

Best Random Forest Params: {'max_depth': 4, 'n_estimators': 50}
Random Forest Grid Search CV Score: 0.9916666666666668

Hyperparameter Tuning for SVM:

Perform hyperparameter tuning for SVM using GridSearchCV.

```
In [162... svm_params = {'C': [0.1, 1, 10], 'gamma': ['scale', 'auto']}
svm_grid = GridSearchCV(SVC(), svm_params, cv=5)
svm_grid.fit(x_train, y_train)
print(f'Best SVM Params: {svm_grid.best_params_}')
print(f'SVM Grid Search CV Score: {svm_grid.best_score_}')
```

Best SVM Params: {'C': 1, 'gamma': 'scale'}
SVM Grid Search CV Score: 0.975

Hyperparameter Tuning for Gradient Boosting

```
In [163... gb_params = {'n_estimators': [100, 200, 300], 'learning_rate': [0.1, 0.01, 0.001]}
gb_grid = GridSearchCV(GradientBoostingClassifier(), gb_params, cv=5)
gb_grid.fit(X_scaled, y)
print(f'Best Gradient Boosting Params: {gb_grid.best_params_}')
print(f'Best Gradient Boosting Accuracy: {gb_grid.best_score_:.2f}')
```

Best Gradient Boosting Params: {'learning_rate': 0.001, 'n_estimators': 100}
Best Gradient Boosting Accuracy: 0.97

In []: